

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier
from xgboost import XGBClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import HistGradientBoostingClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder

from sklearn.model_selection import GridSearchCV

import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: downloaded_dataset = pd.read_csv(
    r"C:\Network Intrusion Detection System\KDDCup Data 10 Percent (1).csv",
    header=None)
print("First 10 Elements from the Dataset - KDDCup Data 10 Percent.csv:")
downloaded_dataset.head(10)
```

First 10 Elements from the Dataset - KDDCup Data 10 Percent.csv:

```
Out[2]:
```

	0	1	2	3	4	5	6	7	8	9	...	32	33	34	35	36	37	38	39	40
0	0	tcp	http	SF	181	5450	0	0	0	0	...	9	1.0	0.0	0.11	0.00	0.0	0.0	0.0	0.0
1	0	tcp	http	SF	239	486	0	0	0	0	...	19	1.0	0.0	0.05	0.00	0.0	0.0	0.0	0.0
2	0	tcp	http	SF	235	1337	0	0	0	0	...	29	1.0	0.0	0.03	0.00	0.0	0.0	0.0	0.0
3	0	tcp	http	SF	219	1337	0	0	0	0	...	39	1.0	0.0	0.03	0.00	0.0	0.0	0.0	0.0
4	0	tcp	http	SF	217	2032	0	0	0	0	...	49	1.0	0.0	0.02	0.00	0.0	0.0	0.0	0.0
5	0	tcp	http	SF	217	2032	0	0	0	0	...	59	1.0	0.0	0.02	0.00	0.0	0.0	0.0	0.0
6	0	tcp	http	SF	212	1940	0	0	0	0	...	69	1.0	0.0	1.00	0.04	0.0	0.0	0.0	0.0
7	0	tcp	http	SF	159	4087	0	0	0	0	...	79	1.0	0.0	0.09	0.04	0.0	0.0	0.0	0.0
8	0	tcp	http	SF	210	151	0	0	0	0	...	89	1.0	0.0	0.12	0.04	0.0	0.0	0.0	0.0
9	0	tcp	http	SF	212	786	0	0	0	1	...	99	1.0	0.0	0.12	0.05	0.0	0.0	0.0	0.0

10 rows × 42 columns



```
In [3]: print("Last 10 Elements from the Dataset - KDDCup Data 10 Percent.csv:")
downloaded_dataset.tail(10)
```

Last 10 Elements from the Dataset - KDDCup Data 10 Percent.csv:

Out[3]:

	0	1	2	3	4	5	6	7	8	9	...	32	33	34	35	36	37	38
494011	0	tcp	http	SF	308	662	0	0	0	0	...	255	1.0	0.0	0.03	0.06	0.00	0.01
494012	0	tcp	http	SF	291	1862	0	0	0	0	...	255	1.0	0.0	0.02	0.05	0.00	0.01
494013	0	tcp	http	SF	289	244	0	0	0	0	...	255	1.0	0.0	0.02	0.05	0.00	0.01
494014	0	tcp	http	SF	306	662	0	0	0	0	...	255	1.0	0.0	0.02	0.05	0.00	0.01
494015	0	tcp	http	SF	289	1862	0	0	0	0	...	255	1.0	0.0	0.01	0.05	0.00	0.01
494016	0	tcp	http	SF	310	1881	0	0	0	0	...	255	1.0	0.0	0.01	0.05	0.00	0.01
494017	0	tcp	http	SF	282	2286	0	0	0	0	...	255	1.0	0.0	0.17	0.05	0.00	0.01
494018	0	tcp	http	SF	203	1200	0	0	0	0	...	255	1.0	0.0	0.06	0.05	0.06	0.01
494019	0	tcp	http	SF	291	1200	0	0	0	0	...	255	1.0	0.0	0.04	0.05	0.04	0.01
494020	0	tcp	http	SF	219	1234	0	0	0	0	...	255	1.0	0.0	0.17	0.05	0.00	0.01

10 rows × 42 columns

In [4]:

```

with open(
    "kddcup.txt", "r"
) as kddcup: # A handle for reading the contents in the File: kddcup.txt
    # Opening the contents from the file, in order to be manipulated...
    kddcup_contents = kddcup.read()
    # I have customized/manipulated the data entry here, by slicing
    data = kddcup_contents[182:].strip("\n")
    # the contents from the starting position of 182; because I only want the column
    # KDDCup Data 10 Percent.csv

# Now we can iterate the contents from the Dataset: KDDCup Data 10
data_scraping = data.split("\n")
# Percent.csv, and store them in the empty List Variable: stored_features.

# for contents in data_scraping:
#     # print(contents) # Un-Comment this to view the contents stored in the List
#     stored_features...
# ----- NOTE -----
# So we know the contents in the List Variable: data_scraping, is an element of a
#
# List Variable: extracted_data - By observance, we know that each Key is the Column
#                               its corresponding connection. This allows us to see
#                               elements are separated as its singular value.
#
# List Variable: extracted_features - Stores the data entry which will be used for
#                               Percent.csv Column Names
#
# List Variable: extracted_connections - Stores the data entry which will be used for
#                               10 Percent.csv connection type.
extracted_data = [
    word for lists_data in data_scraping for word in lists_data.split(": ")

```


Out[7]:

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urg
0	0	tcp	http	SF	181	5450	0	0	
1	0	tcp	http	SF	239	486	0	0	
2	0	tcp	http	SF	235	1337	0	0	
3	0	tcp	http	SF	219	1337	0	0	
4	0	tcp	http	SF	217	2032	0	0	
5	0	tcp	http	SF	217	2032	0	0	
6	0	tcp	http	SF	212	1940	0	0	
7	0	tcp	http	SF	159	4087	0	0	
8	0	tcp	http	SF	210	151	0	0	
9	0	tcp	http	SF	212	786	0	0	

10 rows × 42 columns



In [8]: `print("Last 10 Elements from the Dataset - KDDCup Data 10 Percent (Edited).csv:")`
`edited_dataset[-10:]`

Last 10 Elements from the Dataset - KDDCup Data 10 Percent (Edited).csv:

Out[8]:

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urg
494011	0	tcp	http	SF	308	662	0	0	
494012	0	tcp	http	SF	291	1862	0	0	
494013	0	tcp	http	SF	289	244	0	0	
494014	0	tcp	http	SF	306	662	0	0	
494015	0	tcp	http	SF	289	1862	0	0	
494016	0	tcp	http	SF	310	1881	0	0	
494017	0	tcp	http	SF	282	2286	0	0	
494018	0	tcp	http	SF	203	1200	0	0	
494019	0	tcp	http	SF	291	1200	0	0	
494020	0	tcp	http	SF	219	1234	0	0	

10 rows × 42 columns



The code below will list out all the column names.

In [9]: `print(list(edited_dataset))`

```
['duration', 'protocol_type', 'service', 'flag', 'src_bytes', 'dst_bytes', 'land',
'wrong_fragment', 'urgent', 'hot', 'num_failed_logins', 'logged_in', 'num_compromise
d', 'root_shell', 'su_attempted', 'num_root', 'num_file_creations', 'num_shells', 'n
um_access_files', 'num_outbound_cmds', 'is_host_login', 'is_guest_login', 'count',
'srv_count', 'serror_rate', 'srv_serror_rate', 'rerror_rate', 'srv_rerror_rate', 'sa
me_srv_rate', 'diff_srv_rate', 'srv_diff_host_rate', 'dst_host_count', 'dst_host_srv
_count', 'dst_host_same_srv_rate', 'dst_host_diff_srv_rate', 'dst_host_same_src_port
_rate', 'dst_host_srv_diff_host_rate', 'dst_host_serror_rate', 'dst_host_srv_serror_
_rate', 'dst_host_rerror_rate', 'dst_host_srv_rerror_rate', 'connection']
```

```
In [10]: edited_dataset["connection"].value_counts()
```

```
Out[10]: connection
smurf.          280790
neptune.        107201
normal.         97278
back.           2203
satan.          1589
ipsweep.        1247
portsweep.      1040
warezclient.    1020
teardrop.       979
pod.            264
nmap.           231
guess_passwd.   53
buffer_overflow. 30
land.           21
warezmaster.    20
imap.           12
rootkit.        10
loadmodule.     9
ftp_write.      8
multihop.       7
phf.            4
perl.           3
spy.            2
Name: count, dtype: int64
```

The cell above displays counts of each type of connections.

Let's store all the records in a dataframe normal where connection is 'normal'.

```
In [11]: normal = edited_dataset[edited_dataset.connection == "normal."]
normal.head()
```

Out[11]:

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urg
0	0	tcp	http	SF	181	5450	0	0	
1	0	tcp	http	SF	239	486	0	0	
2	0	tcp	http	SF	235	1337	0	0	
3	0	tcp	http	SF	219	1337	0	0	
4	0	tcp	http	SF	217	2032	0	0	

5 rows × 42 columns

```
In [12]: normal_rows = normal.shape[0]
normal_columns = normal.shape[1]
edited_dataset_rows = edited_dataset.shape[0]
edited_dataset_columns = edited_dataset.shape[1]

print(
    f"The Dataset - normal has {normal_rows:,} rows and {normal_columns} columns."
    + "\n"
    + f"The Dataset - edited_dataset has {edited_dataset_rows:,} rows and {edited_d"
)
```

The Dataset - normal has 97,278 rows and 42 columns.

The Dataset - edited_dataset has 494,021 rows and 42 columns.

Prepare your personalised dataset for this task. Replace '26577' with the last 5 digits of your student number. If it starts with zero, please use 6 as the first digit.

Save the data in to a csv file using the codes in the cells below.

```
In [13]: rows = np.random.choice(edited_dataset.index.values, 24861)
dataset = edited_dataset.loc[rows]

dataset.to_csv("Latest.csv", index=False)
```

Replace name and number with your name and student number. Write it to a csv file. You are required to use this dataset to complete this task.

```
In [14]: dataset = pd.read_csv("Latest.csv")

print("First 10 Elements from the Dataset - KDDCup Data 10 Percent (Edited).csv:")
dataset[:10]
```

First 10 Elements from the Dataset - KDDCup Data 10 Percent (Edited).csv:

Out[14]:

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urg
0	0	tcp	http	SF	341	16123	0	0	
1	2638	udp	other	SF	145	105	0	0	
2	0	icmp	ecr_i	SF	1032	0	0	0	
3	0	tcp	private	REJ	0	0	0	0	
4	0	icmp	ecr_i	SF	1032	0	0	0	
5	0	tcp	private	S0	0	0	0	0	
6	0	tcp	http	SF	320	334	0	0	
7	0	icmp	ecr_i	SF	520	0	0	0	
8	0	icmp	ecr_i	SF	1032	0	0	0	
9	0	tcp	smtp	SF	1105	328	0	0	

10 rows × 42 columns



In [15]: `print("Last 10 Elements from the Dataset - KDDCup Data 10 Percent (Edited).csv:")`
`dataset[-10:]`

Last 10 Elements from the Dataset - KDDCup Data 10 Percent (Edited).csv:

Out[15]:

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment
24851	0	tcp	http	SF	345	398	0	0
24852	0	tcp	http	SF	343	6177	0	0
24853	0	icmp	ecr_i	SF	1032	0	0	0
24854	0	icmp	ecr_i	SF	1032	0	0	0
24855	0	icmp	ecr_i	SF	1032	0	0	0
24856	0	tcp	private	S0	0	0	0	0
24857	0	tcp	private	S0	0	0	0	0
24858	4	tcp	smtp	SF	699	328	0	0
24859	0	icmp	ecr_i	SF	1032	0	0	0
24860	0	icmp	ecr_i	SF	1032	0	0	0

10 rows × 42 columns



In [16]: `dataset["connection"].value_counts()`

```
Out[16]: connection
smurf.          14066
neptune.        5429
normal.         4905
back.           110
ipsweep.        76
satan.          64
warezclient.    62
portsweep.      57
teardrop.       55
pod.            14
nmap.           12
guess_passwd.   5
buffer_overflow. 3
loadmodule.     2
imap.           1
Name: count, dtype: int64
```

```
In [17]: dataset.head()
```

```
Out[17]:
```

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urg
0	0	tcp	http	SF	341	16123	0	0	
1	2638	udp	other	SF	145	105	0	0	
2	0	icmp	ecr_i	SF	1032	0	0	0	
3	0	tcp	private	REJ	0	0	0	0	
4	0	icmp	ecr_i	SF	1032	0	0	0	

5 rows × 42 columns



Data Preprocessing

```
In [18]: # will return the data types of all the columns
dataset.dtypes
```



```

Out[18]: duration                int64
         protocol_type           object
         service                 object
         flag                   object
         src_bytes               int64
         dst_bytes               int64
         land                    int64
         wrong_fragment          int64
         urgent                  int64
         hot                     int64
         num_failed_logins        int64
         logged_in               int64
         num_compromised          int64
         root_shell              int64
         su_attempted            int64
         num_root                 int64
         num_file_creations       int64
         num_shells               int64
         num_access_files         int64
         num_outbound_cmds        int64
         is_host_login            int64
         is_guest_login           int64
         count                    int64
         srv_count                int64
         serror_rate              float64
         srv_serror_rate          float64
         rerror_rate              float64
         srv_rerror_rate          float64
         same_srv_rate            float64
         diff_srv_rate            float64
         srv_diff_host_rate       float64
         dst_host_count           int64
         dst_host_srv_count       int64
         dst_host_same_srv_rate   float64
         dst_host_diff_srv_rate   float64
         dst_host_same_src_port_rate float64
         dst_host_srv_diff_host_rate float64
         dst_host_serror_rate     float64
         dst_host_srv_serror_rate float64
         dst_host_rerror_rate     float64
         dst_host_srv_rerror_rate float64
         connection               object
         dtype: object

```

```
In [19]: dataset.shape
```

```
Out[19]: (24861, 42)
```

```

In [20]: # will generate descriptive statistics of your dataset
         dataset.describe()

```

Out[20]:

	duration	src_bytes	dst_bytes	land	wrong_fragment	urgent	
count	24861.000000	2.486100e+04	24861.000000	24861.0	24861.000000	24861.0	2486
mean	58.158039	1.878291e+03	606.907244	0.0	0.006999	0.0	
std	794.623408	6.811794e+04	4060.136078	0.0	0.140506	0.0	
min	0.000000	0.000000e+00	0.000000	0.0	0.000000	0.0	
25%	0.000000	4.300000e+01	0.000000	0.0	0.000000	0.0	
50%	0.000000	5.200000e+02	0.000000	0.0	0.000000	0.0	
75%	0.000000	1.032000e+03	0.000000	0.0	0.000000	0.0	
max	31061.000000	5.135678e+06	288690.000000	0.0	3.000000	0.0	3

8 rows × 38 columns



In [21]: dataset.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24861 entries, 0 to 24860
Data columns (total 42 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   duration                             24861 non-null  int64
 1   protocol_type                        24861 non-null  object
 2   service                              24861 non-null  object
 3   flag                                 24861 non-null  object
 4   src_bytes                           24861 non-null  int64
 5   dst_bytes                           24861 non-null  int64
 6   land                                24861 non-null  int64
 7   wrong_fragment                      24861 non-null  int64
 8   urgent                             24861 non-null  int64
 9   hot                                 24861 non-null  int64
10  num_failed_logins                   24861 non-null  int64
11  logged_in                           24861 non-null  int64
12  num_compromised                     24861 non-null  int64
13  root_shell                          24861 non-null  int64
14  su_attempted                       24861 non-null  int64
15  num_root                            24861 non-null  int64
16  num_file_creations                 24861 non-null  int64
17  num_shells                         24861 non-null  int64
18  num_access_files                   24861 non-null  int64
19  num_outbound_cmds                 24861 non-null  int64
20  is_host_login                      24861 non-null  int64
21  is_guest_login                     24861 non-null  int64
22  count                             24861 non-null  int64
23  srv_count                         24861 non-null  int64
24  serror_rate                        24861 non-null  float64
25  srv_serror_rate                    24861 non-null  float64
26  rerror_rate                        24861 non-null  float64
27  srv_rerror_rate                    24861 non-null  float64
28  same_srv_rate                      24861 non-null  float64
29  diff_srv_rate                      24861 non-null  float64
30  srv_diff_host_rate                 24861 non-null  float64
31  dst_host_count                     24861 non-null  int64
32  dst_host_srv_count                 24861 non-null  int64
33  dst_host_same_srv_rate             24861 non-null  float64
34  dst_host_diff_srv_rate             24861 non-null  float64
35  dst_host_same_src_port_rate        24861 non-null  float64
36  dst_host_srv_diff_host_rate        24861 non-null  float64
37  dst_host_serror_rate               24861 non-null  float64
38  dst_host_srv_serror_rate           24861 non-null  float64
39  dst_host_rerror_rate               24861 non-null  float64
40  dst_host_srv_rerror_rate           24861 non-null  float64
41  connection                         24861 non-null  object
dtypes: float64(15), int64(23), object(4)
memory usage: 8.0+ MB

```

```

In [22]: # Tells how many missing values are present in each column
dataset.isnull().sum()

```

```

Out[22]: duration          0
         protocol_type     0
         service           0
         flag              0
         src_bytes         0
         dst_bytes         0
         land              0
         wrong_fragment    0
         urgent            0
         hot               0
         num_failed_logins  0
         logged_in         0
         num_compromised   0
         root_shell        0
         su_attempted      0
         num_root          0
         num_file_creations 0
         num_shells        0
         num_access_files  0
         num_outbound_cmds 0
         is_host_login     0
         is_guest_login    0
         count             0
         srv_count         0
         serror_rate       0
         srv_serror_rate   0
         rerror_rate       0
         srv_rerror_rate   0
         same_srv_rate     0
         diff_srv_rate     0
         srv_diff_host_rate 0
         dst_host_count    0
         dst_host_srv_count 0
         dst_host_same_srv_rate 0
         dst_host_diff_srv_rate 0
         dst_host_same_src_port_rate 0
         dst_host_srv_diff_host_rate 0
         dst_host_serror_rate 0
         dst_host_srv_serror_rate 0
         dst_host_rerror_rate 0
         dst_host_srv_rerror_rate 0
         connection       0
         dtype: int64

```

```

In [23]: dataset.drop_duplicates(inplace=True)
         dataset.shape

```

```

Out[23]: (9807, 42)

```

```

In [24]: dataset.columns

```

```
Out[24]: Index(['duration', 'protocol_type', 'service', 'flag', 'src_bytes',
               'dst_bytes', 'land', 'wrong_fragment', 'urgent', 'hot',
               'num_failed_logins', 'logged_in', 'num_compromised', 'root_shell',
               'su_attempted', 'num_root', 'num_file_creations', 'num_shells',
               'num_access_files', 'num_outbound_cmds', 'is_host_login',
               'is_guest_login', 'count', 'srv_count', 'serror_rate',
               'srv_serror_rate', 'rerror_rate', 'srv_rerror_rate', 'same_srv_rate',
               'diff_srv_rate', 'srv_diff_host_rate', 'dst_host_count',
               'dst_host_srv_count', 'dst_host_same_srv_rate',
               'dst_host_diff_srv_rate', 'dst_host_same_src_port_rate',
               'dst_host_srv_diff_host_rate', 'dst_host_serror_rate',
               'dst_host_srv_serror_rate', 'dst_host_rerror_rate',
               'dst_host_srv_rerror_rate', 'connection'],
              dtype='object')
```

```
In [25]: # This code is used to find all categorical (string/object) columns in a Pandas Dat

obj_cols = dataset.select_dtypes(include=["object"]).columns
print(obj_cols)
```

```
Index(['protocol_type', 'service', 'flag', 'connection'], dtype='object')
```

```
In [26]: #This code is used to convert categorical (text) columns into numerical values so t

le = LabelEncoder()
for i in obj_cols:
    dataset[i] = le.fit_transform(dataset[i])
```

```
In [27]: dataset.dtypes
```

```

Out[27]: duration                int64
         protocol_type           int64
         service                 int64
         flag                    int64
         src_bytes               int64
         dst_bytes               int64
         land                    int64
         wrong_fragment          int64
         urgent                  int64
         hot                     int64
         num_failed_logins       int64
         logged_in               int64
         num_compromised         int64
         root_shell              int64
         su_attempted            int64
         num_root                int64
         num_file_creations      int64
         num_shells              int64
         num_access_files        int64
         num_outbound_cmds       int64
         is_host_login           int64
         is_guest_login          int64
         count                   int64
         srv_count               int64
         serror_rate             float64
         srv_serror_rate         float64
         rerror_rate             float64
         srv_rerror_rate         float64
         same_srv_rate           float64
         diff_srv_rate           float64
         srv_diff_host_rate      float64
         dst_host_count          int64
         dst_host_srv_count      int64
         dst_host_same_srv_rate  float64
         dst_host_diff_srv_rate  float64
         dst_host_same_src_port_rate float64
         dst_host_srv_diff_host_rate float64
         dst_host_serror_rate    float64
         dst_host_srv_serror_rate float64
         dst_host_rerror_rate    float64
         dst_host_srv_rerror_rate float64
         connection              int64
         dtype: object

```

In [28]: *#This code is used to get all column names except the last column.*

```

columns = list(dataset)
columns = columns[:-1]

```

In [29]: *#This code is used to standardize (scale) the numerical features before training a
StandardScaler standardizes numerical features so that each feature has a mean of*

```

scaler = StandardScaler()
sc_data = scaler.fit_transform(dataset[columns])
print(sc_data)

```

```

[[-0.11174437 -0.1907899 -1.01381963 ... -0.79526687 -0.39813978
  -0.39389014]
 [ 2.00597162  2.82707544  0.45846906 ... -0.79526687 -0.39813978
  -0.39389014]
 [-0.11174437 -3.20865523 -1.66817015 ... -0.79526687 -0.39813978
  -0.39389014]
 ...
 [-0.11174437 -0.1907899 -1.01381963 ... -0.79526687 -0.36826902
  -0.3639093 ]
 [-0.11174437 -0.1907899  0.78564432 ...  1.26033027 -0.39813978
  -0.39389014]
 [-0.10853328 -0.1907899  1.11281959 ... -0.79526687 -0.39813978
  -0.39389014]]

```

```

In [30]: # TRANSFORMED DATASET
transformed_data = pd.DataFrame(sc_data, columns=columns)
transformed_data.head()

```

```

Out[30]:
   duration  protocol_type  service  flag  src_bytes  dst_bytes  land  wrong_fragment
0 -0.111744    -0.190790 -1.013820  0.882304 -0.026956  2.319289  0.0    -0.0777
1  2.005972     2.827075  0.458469  0.882304 -0.028764 -0.217609  0.0    -0.0777
2 -0.111744    -3.208655 -1.668170  0.882304 -0.020582 -0.234239  0.0    -0.0777
3 -0.111744    -0.190790  0.785644 -2.102140 -0.030102 -0.234239  0.0    -0.0777
4 -0.111744    -0.190790  0.785644 -0.396743 -0.030102 -0.234239  0.0    -0.0777

```

5 rows × 41 columns



```

In [31]: full = pd.DataFrame(transformed_data)
full

```

Out[31]:

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fra
0	-0.111744	-0.190790	-1.013820	0.882304	-0.026956	2.319289	0.0	-0.0
1	2.005972	2.827075	0.458469	0.882304	-0.028764	-0.217609	0.0	-0.0
2	-0.111744	-3.208655	-1.668170	0.882304	-0.020582	-0.234239	0.0	-0.0
3	-0.111744	-0.190790	0.785644	-2.102140	-0.030102	-0.234239	0.0	-0.0
4	-0.111744	-0.190790	0.785644	-0.396743	-0.030102	-0.234239	0.0	-0.0
...	...	...	...	...	...	...	...	...
9802	-0.111744	-0.190790	0.785644	-0.396743	-0.030102	-0.234239	0.0	-0.0
9803	-0.111744	-0.190790	-1.013820	0.882304	-0.026919	-0.171204	0.0	-0.0
9804	-0.111744	-0.190790	-1.013820	0.882304	-0.026938	0.744062	0.0	-0.0
9805	-0.111744	-0.190790	0.785644	-0.396743	-0.030102	-0.234239	0.0	-0.0
9806	-0.108533	-0.190790	1.112820	0.882304	-0.023654	-0.182291	0.0	-0.0

9807 rows × 41 columns



In [32]: *#This creates a new column called connection inside full.*
 full["connection"] = dataset["connection"]
 full.head()

Out[32]:

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment
0	-0.111744	-0.190790	-1.013820	0.882304	-0.026956	2.319289	0.0	-0.0777
1	2.005972	2.827075	0.458469	0.882304	-0.028764	-0.217609	0.0	-0.0777
2	-0.111744	-3.208655	-1.668170	0.882304	-0.020582	-0.234239	0.0	-0.0777
3	-0.111744	-0.190790	0.785644	-2.102140	-0.030102	-0.234239	0.0	-0.0777
4	-0.111744	-0.190790	0.785644	-0.396743	-0.030102	-0.234239	0.0	-0.0777

5 rows × 42 columns




In [33]: *# connection → Target (Dependent variable)*
 target = dataset["connection"]
 target.head(10)


```
Out[33]: 0      8
         1      8
         2     12
         3      6
         4      6
         5      8
         6     12
         7     12
         8      8
         9     12
        10     12
        Name: connection, dtype: int64
```

Insights about the Data


```
In [34]: # rcParams stands for Runtime Configuration Parameters.
         # It is a dictionary that stores Matplotlib's default settings, such as:
         plt.rcParams["figure.figsize"] = (12, 8)
```

```
In [35]: dataset['protocol_type'].value_counts().plot(kind='bar')
         plt.xlabel('Protocol Type')
         plt.ylabel('Count')
         plt.title('Protocol Type Distribution')
         plt.show()
```




```
In [36]: dataset['duration'].plot(kind='hist', bins=20)
         plt.xlabel('Duration')
         plt.ylabel('Frequency')
```


```
plt.title('Duration Distribution')  
plt.show()
```




```
In [37]: plt.scatter(dataset['src_bytes'], dataset['dst_bytes'])  
plt.xlabel('Source Bytes')  
plt.ylabel('Destination Bytes')  
plt.title('Source Bytes vs Destination Bytes')  
plt.show()
```




```
In [38]: dataset.boxplot(column='error_rate')  
plt.ylabel('Error Rate')  
plt.title('Error Rate Distribution')  
plt.show()
```




```
In [39]: dataset['flag'].value_counts().plot(kind='bar')
plt.xlabel('Flag')
plt.ylabel('Count')
plt.title('Flag Distribution')
plt.show()
```




```
In [40]: dataset['hot'].plot(kind='hist', bins=20)
plt.xlabel('Hot')
plt.ylabel('Frequency')
plt.title('Hot Distribution')
plt.show()
```




```
In [41]: dataset['logged_in'].value_counts().plot(kind='bar')
plt.xlabel('Logged In')
plt.ylabel('Count')
plt.title('Logged In Distribution')
plt.show()
```




```
In [42]: dataset.boxplot(column='num_failed_logins')
plt.ylabel('Number of Failed Logins')
plt.title('Number of Failed Logins Distribution')
plt.show()
```




```
In [43]: dataset['is_host_login'].value_counts().plot(kind='bar')
plt.xlabel('Is Host Login')
plt.ylabel('Count')
plt.title('Is Host Login Distribution')
plt.show()
```




```
In [44]: dataset['is_guest_login'].value_counts().plot(kind='bar')
plt.xlabel('Is Guest Login')
plt.ylabel('Count')
plt.title('Is Guest Login Distribution')
plt.show()
```




```
In [45]: dataset['count'].plot(kind='hist', bins=20)
plt.xlabel('Count')
plt.ylabel('Frequency')
plt.title('Count Distribution')
plt.show()
```




```
In [46]: dataset['error_rate'].plot(kind='hist', bins=20)
plt.xlabel('Error Rate')
plt.ylabel('Frequency')
plt.title('Error Rate Distribution')
plt.show()
```


```
In [47]: plt.scatter(dataset['srv_error_rate'], dataset['rerror_rate'])
plt.xlabel('Srv Error Rate')
plt.ylabel('Rerror Rate')
plt.title('Srv Error Rate vs Rerror Rate')
plt.show()
```




```
In [48]: dataset.boxplot(column='dst_host_same_srv_rate')  
plt.ylabel('Dst Host Same Srv Rate')  
plt.title('Dst Host Same Srv Rate Distribution')  
plt.show()
```




```
In [49]: dataset.boxplot(column='dst_host_diff_srv_rate')  
plt.ylabel('Dst Host Diff Srv Rate')  
plt.title('Dst Host Diff Srv Rate Distribution')  
plt.show()
```



```
In [50]: plt.scatter(dataset['dst_host_srv_count'], dataset['dst_host_same_srv_rate'])  
plt.xlabel('Dst Host Srv Count')  
plt.ylabel('Dst Host Same Srv Rate')  
plt.title('Dst Host Srv Count vs Dst Host Same Srv Rate')  
plt.show()
```




```
In [51]: dataset['connection'].value_counts().plot(kind='bar')
plt.xlabel('Connection')
plt.ylabel('Count')
plt.title('Connection Distribution')
plt.show()
```



```
In [52]: corr = dataset.corr()
plt.figure(figsize=(20, 20))
plt.title('Heatmap of Correlation Matrix')
sns.heatmap(corr, annot=True, cmap='coolwarm')
```

```
Out[52]: <Axes: title={'center': 'Heatmap of Correlation Matrix'}>
```



Training and Splitting Data set into Test and Train

```
In [53]: # Splitting the dataset into the Training set and Test set:
X_train, X_test, y_train, y_test = train_test_split(
    transformed_data, target, test_size=0.3, random_state=8)
```

Different Classification Models

Decision Tree Classifier

```
In [54]: from sklearn.tree import DecisionTreeClassifier

decision_tree = DecisionTreeClassifier()

parameters = [
    {
        "criterion": ["gini", "entropy"],
        "splitter": ["best", "random"],
        "max_features": ["auto", "sqrt", "log2", None],
    }
]

random_search = GridSearchCV(decision_tree, parameters)
random_search.fit(X_train, y_train)
random_search.best_params_
```

```
Out[54]: {'criterion': 'gini', 'max_features': None, 'splitter': 'random'}
```

```
In [55]: print("Evaluation for Decision Tree Classifier".center(75, '_'))

decision_tree = DecisionTreeClassifier()
decision_tree.fit(X_train, y_train)

decision_tree_prediction = decision_tree.predict(X_test)
decision_tree_accuracy = accuracy_score(y_test, decision_tree_prediction)
decision_tree_precision = precision_score(
    y_test, decision_tree_prediction, average="weighted")
decision_tree_recall = recall_score(
    y_test, decision_tree_prediction, average="weighted")
decision_tree_f1_score = f1_score(
    y_test, decision_tree_prediction, average="weighted")

print("Prediciton: ", decision_tree_prediction)
print('_' * 75)

print("Accuracy:" + "\t" + f"{{(decision_tree_accuracy * 100)}}%")
print("Precision:" + "\t" + f"{{(decision_tree_precision * 100)}}%")
print("Recall:" + "\t\t" + f"{{(decision_tree_recall * 100)}}%")
print("F1-Score:" + "\t" + f"{{(decision_tree_f1_score * 100)}}%")
print('_' * 75)
```

```
_____Evaluation for Decision Tree Classifier_____
Prediciton: [6 8 6 ... 6 8 6]

Accuracy: 99.28644240570847%
Precision: 99.36386572766747%
Recall: 99.28644240570847%
F1-Score: 99.29126845327966%
```

```
In [56]: print("Confusion Matrix For Decision Tree Classifier - 'Labels Test' & 'Prediction'")
print("-" * 93)
print(confusion_matrix(y_test, decision_tree_prediction))
```

Confusion Matrix For Decision Tree Classifier - 'Labels Test' & 'Prediction':

[	28	0	0	0	0	0	0	0	0	0	0	0	0]
[	0	1	0	0	0	0	0	0	0	0	0	0	0]
[	0	0	2	0	0	0	0	0	0	0	0	0	0]
[	0	0	0	0	0	1	0	0	0	0	0	0	0]
[	0	0	0	0	22	0	0	0	0	0	0	0	0]
[	0	0	0	0	0	1400	0	0	0	0	0	0	0]
[	0	0	0	0	1	0	0	0	0	0	0	0	0]
[	0	0	0	0	1	2	0	1387	0	0	2	1	0]
[	0	0	0	0	0	0	0	4	0	0	0	0	0]
[	0	0	0	0	2	0	1	0	0	13	2	0	0]
[	0	0	0	0	0	0	0	0	0	15	0	0	0]
[	0	0	0	0	0	0	1	0	0	0	27	0	0]
[	0	0	0	0	0	0	0	0	0	0	0	10	0]
[	0	0	0	0	0	0	2	0	0	0	0	0	13]]

```
In [57]: print(
          "Classification Report For Decision Tree Classifier - 'Labels Test' & 'Prediction'
          )
          print("-" * 97)
          print(classification_report(y_test, decision_tree_prediction))
```

Classification Report For Decision Tree Classifier - 'Labels Test' & 'Prediction':


	precision	recall	f1-score	support
0	1.00	1.00	1.00	28
1	1.00	1.00	1.00	1
2	1.00	1.00	1.00	2
3	0.00	0.00	0.00	1
4	0.85	1.00	0.92	22
6	1.00	1.00	1.00	1400
7	0.00	0.00	0.00	1
8	1.00	0.99	1.00	1397
9	1.00	1.00	1.00	4
10	1.00	0.68	0.81	19
11	0.79	1.00	0.88	15
12	0.96	0.96	0.96	28
13	1.00	1.00	1.00	10
14	0.72	0.87	0.79	15
accuracy			0.99	2943
macro avg	0.81	0.82	0.81	2943
weighted avg	0.99	0.99	0.99	2943

```
In [58]: feature_importance = pd.DataFrame(decision_tree.feature_importances_, index=transformer.get_feature_names_out(),
          'importance']).sort_values('importance', ascending=False)
          print(feature_importance)
```


	importance
same_srv_rate	0.798909
src_bytes	0.042739
srv_count	0.037928
diff_srv_rate	0.024470
wrong_fragment	0.021005
dst_host_srv_diff_host_rate	0.018326
serror_rate	0.015743
service	0.010329
dst_host_same_src_port_rate	0.007835
hot	0.005156
dst_host_diff_srv_rate	0.003682
dst_host_srv_count	0.002446
flag	0.001901
srv_rerror_rate	0.001512
protocol_type	0.001445
dst_host_srv_rerror_rate	0.001218
rerror_rate	0.001150
root_shell	0.001056
dst_bytes	0.000877
dst_host_same_srv_rate	0.000873
num_failed_logins	0.000703
is_guest_login	0.000521
dst_host_serror_rate	0.000177
duration	0.000000
land	0.000000
urgent	0.000000
num_access_files	0.000000
count	0.000000
is_host_login	0.000000
num_shells	0.000000
su_attempted	0.000000
num_root	0.000000
num_file_creations	0.000000
logged_in	0.000000
num_compromised	0.000000
num_outbound_cmds	0.000000
srv_serror_rate	0.000000
dst_host_count	0.000000
srv_diff_host_rate	0.000000
dst_host_srv_serror_rate	0.000000
dst_host_rerror_rate	0.000000

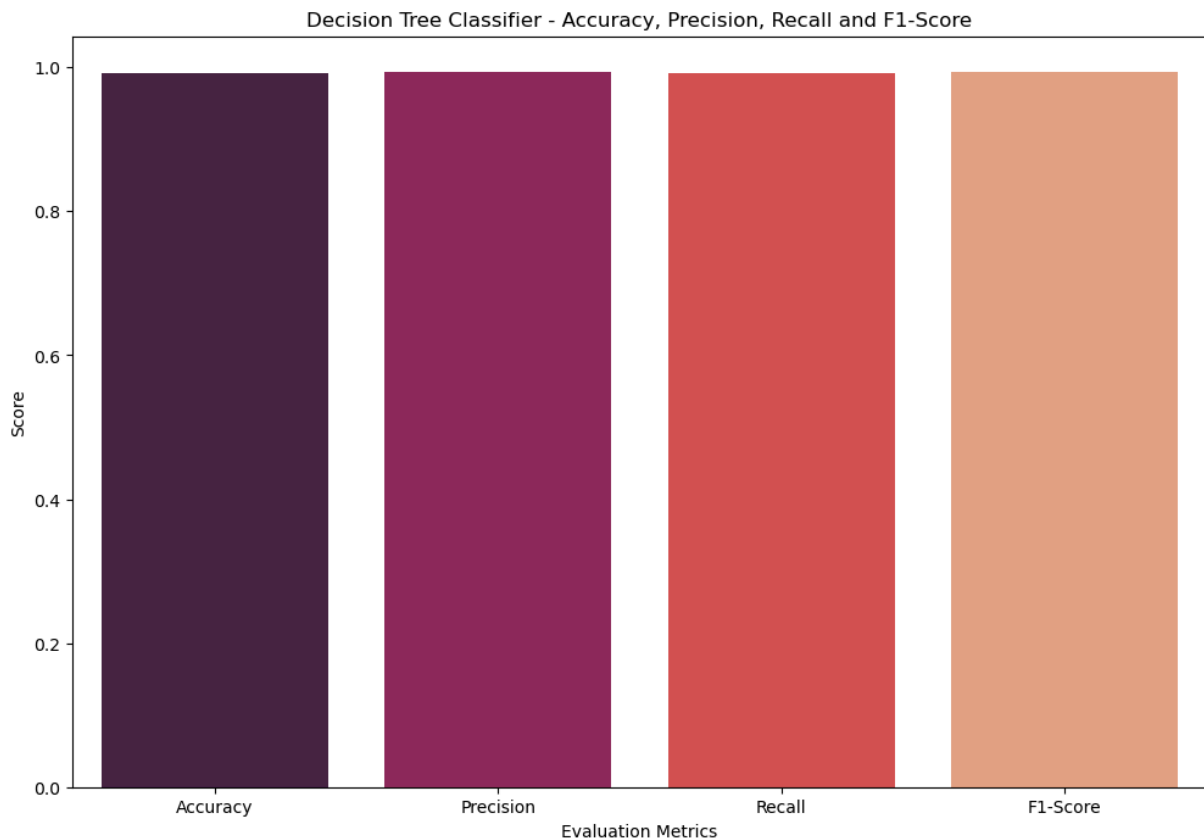
```
In [59]: feature_importance.plot(kind="barh")
```

```
Out[59]: <Axes: >
```



```
In [60]: season = ["Accuracy", "Precision", "Recall", "F1-Score"]
season_score = [decision_tree_accuracy, decision_tree_precision,
                decision_tree_recall, decision_tree_f1_score]

sns.barplot(x=season, y=season_score, palette="rocket")
plt.title("Decision Tree Classifier - Accuracy, Precision, Recall and F1-Score")
plt.xlabel("Evaluation Metrics")
plt.ylabel("Score")
plt.show()
```



Random Forest Classifier

```
In [61]: from sklearn.ensemble import RandomForestClassifier

random_forest = RandomForestClassifier()

parameters = [
    {
        "n_estimators": [10, 50, 100, 150, 200],
        "criterion": ["gini", "entropy"],
        "max_features": ["auto", "sqrt", "log2", None],
    }
]

random_search = GridSearchCV(random_forest, parameters)
random_search.fit(X_train, y_train)
random_search.best_params_
```

```
Out[61]: {'criterion': 'gini', 'max_features': 'sqrt', 'n_estimators': 150}
```

```
In [62]: print("Evaluation for Random Forest Classifier".center(75, '_'))

random_forest_classifier = RandomForestClassifier(
    n_estimators=100, random_state=0)
random_forest_classifier.fit(X_train, y_train)

random_forest_classifier_prediction = random_forest_classifier.predict(X_test)
random_forest_classifier_accuracy = accuracy_score(
```

```

y_test, random_forest_classifier_prediction)
random_forest_classifier_precision = precision_score(
    y_test, random_forest_classifier_prediction, average="weighted")
random_forest_classifier_recall = recall_score(
    y_test, random_forest_classifier_prediction, average="weighted")
random_forest_classifier_f1_score = f1_score(
    y_test, random_forest_classifier_prediction, average="weighted")

print("Prediciton:    ", random_forest_classifier_prediction)
print('_', * 75)

print("Accuracy:" + "\t" + f"{(random_forest_classifier_accuracy * 100)}%")
print("Precision:" + "\t" + f"{(random_forest_classifier_precision * 100)}%")
print("Recall:" + "\t\t" + f"{(random_forest_classifier_recall * 100)}%")
print("F1-Score:" + "\t" + f"{(random_forest_classifier_f1_score * 100)}%")
print('_', * 75)

```

Evaluation for Random Forest Classifier

Prediciton: [6 8 6 ... 6 8 6]

Accuracy:	99.79612640163099%
Precision:	99.7339901912306%
Recall:	99.79612640163099%
F1-Score:	99.75609448621405%

```

In [63]: print("Confusion Matrix For Random Forest Classifier - 'Labels Test' & 'Prediction'")
print("-" * 93)
print(confusion_matrix(y_test, random_forest_classifier_prediction))

```

Confusion Matrix For Random Forest Classifier - 'Labels Test' & 'Prediction':

```

-----
-----
[[ 28  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  1  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  1  0  0  0  0  1  0  0  0  0  0  0]
 [ 0  0  0  0  0  1  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  22  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0 1400  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  1  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0 1396  0  0  0  0  0  1]
 [ 0  0  0  0  0  0  0  0  4  0  0  0  0  0]
 [ 0  0  0  0  0  1  0  0  0  17  1  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  15  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  28  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  10  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0 15]]

```

```

In [64]: print(
    "Classification Report For Random Forest Classifier - 'Labels Test' & 'Predicti
    ")
print("-" * 97)
print(classification_report(y_test, random_forest_classifier_prediction))

```

Classification Report For Random Forest Classifier - 'Labels Test' & 'Prediction':

	precision	recall	f1-score	support
0	1.00	1.00	1.00	28
1	1.00	1.00	1.00	1
2	1.00	0.50	0.67	2
3	0.00	0.00	0.00	1
4	0.96	1.00	0.98	22
6	1.00	1.00	1.00	1400
7	0.00	0.00	0.00	1
8	1.00	1.00	1.00	1397
9	1.00	1.00	1.00	4
10	1.00	0.89	0.94	19
11	0.94	1.00	0.97	15
12	1.00	1.00	1.00	28
13	1.00	1.00	1.00	10
14	0.94	1.00	0.97	15
accuracy			1.00	2943
macro avg	0.84	0.81	0.82	2943
weighted avg	1.00	1.00	1.00	2943

```
In [65]: feature_importance = pd.Series(
    random_forest_classifier.feature_importances_, index=transformed_data.columns)
feature_importance
```


```

Out[65]: duration                2.387013e-03
         protocol_type           1.255539e-02
         service                 1.043770e-02
         flag                    8.746110e-02
         src_bytes               9.899253e-02
         dst_bytes              1.996021e-02
         land                   0.000000e+00
         wrong_fragment          1.143382e-02
         urgent                  0.000000e+00
         hot                    6.585637e-03
         num_failed_logins       1.072640e-04
         logged_in              1.168597e-02
         num_compromised        8.612565e-03
         root_shell             1.560736e-04
         su_attempted           0.000000e+00
         num_root               4.979625e-05
         num_file_creations      2.682452e-05
         num_shells             1.781059e-09
         num_access_files       4.978477e-10
         num_outbound_cmds      0.000000e+00
         is_host_login          0.000000e+00
         is_guest_login         9.206832e-04
         count                  1.473580e-01
         srv_count              1.540508e-02
         serror_rate            2.694794e-02
         srv_serror_rate        1.954778e-02
         rerror_rate            8.880578e-03
         srv_rerror_rate        4.903293e-03
         same_srv_rate          1.908373e-01
         diff_srv_rate          8.688497e-02
         srv_diff_host_rate     1.280935e-03
         dst_host_count         1.595450e-02
         dst_host_srv_count     3.298203e-02
         dst_host_same_srv_rate 3.152031e-02
         dst_host_diff_srv_rate 3.172841e-02
         dst_host_same_src_port_rate 2.170974e-02
         dst_host_srv_diff_host_rate 1.595320e-02
         dst_host_serror_rate   4.822595e-02
         dst_host_srv_serror_rate 1.860171e-02
         dst_host_rerror_rate   4.082593e-03
         dst_host_srv_rerror_rate 5.823023e-03
         dtype: float64

```


```
In [66]: feature_importance.plot(kind="barh")
```

```
Out[66]: <Axes: >
```



```
In [67]: season = ["Accuracy", "Precision", "Recall", "F1-Score"]
season_score = [random_forest_classifier_accuracy, random_forest_classifier_precision,
                random_forest_classifier_recall, random_forest_classifier_f1_score]

sns.barplot(x=season, y=season_score, palette="rocket")
plt.title("Random Forest Classifier - Accuracy, Precision, Recall and F1-Score")
plt.xlabel("Evaluation Metrics")
plt.ylabel("Score")
plt.show()
```



K-Nearest Neighbors

```
In [68]: # apply grid search to knn
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier

knn_model = KNeighborsClassifier()

knn_param_grid = [
    {
        'n_neighbors': [3, 5, 7, 9, 11, 13, 15],
        'weights': ['uniform', 'distance'],
        'metric': ['euclidean', 'manhattan']
    }
]

knn_grid_search = GridSearchCV(knn_model, knn_param_grid)
knn_grid_search.fit(X_train, y_train)
knn_grid_search.best_params_
```

```
Out[68]: {'metric': 'euclidean', 'n_neighbors': 3, 'weights': 'distance'}
```

```
In [69]: print("Evaluation for K-Nearest Neighbors".center(75, '_'))

k_nearest_neighbors = KNeighborsClassifier(n_neighbors=5)
k_nearest_neighbors.fit(X_train, y_train)

k_nearest_neighbors_prediction = k_nearest_neighbors.predict(X_test)
k_nearest_neighbors_accuracy = accuracy_score(
    y_test, k_nearest_neighbors_prediction)
k_nearest_neighbors_precision = precision_score(
    y_test, k_nearest_neighbors_prediction, average="weighted")
k_nearest_neighbors_recall = recall_score(
    y_test, k_nearest_neighbors_prediction, average="weighted")
k_nearest_neighbors_f1_score = f1_score(
    y_test, k_nearest_neighbors_prediction, average="weighted")

print("Prediciton:      ", k_nearest_neighbors_prediction)
print('_' * 75)

print("Accuracy:" + "\t" + f"{(k_nearest_neighbors_accuracy * 100)}%")
print("Precision:" + "\t" + f"{(k_nearest_neighbors_precision * 100)}%")
print("Recall:" + "\t\t" + f"{(k_nearest_neighbors_recall * 100)}%")
print("F1-Score:" + "\t" + f"{(k_nearest_neighbors_f1_score * 100)}%")
print('_' * 75)
```

Evaluation for K-Nearest Neighbors	
Prediciton:	[6 8 6 ... 6 8 6]
Accuracy:	99.11654774040095%
Precision:	98.99119621093692%
Recall:	99.11654774040095%
F1-Score:	98.9979778342897%


```
In [70]: print("Confusion Matrix For K-Nearest Neighbors - 'Labels Test' & 'Prediction':")
print('-' * 93)
print(confusion_matrix(y_test, k_nearest_neighbors_prediction))
```

Confusion Matrix For K-Nearest Neighbors - 'Labels Test' & 'Prediction':

```
-----
[[ 27  0  0  0  0  0  0  1  0  0  0  0  0  0]
 [ 1  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  2  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  1  0  0  0  0  0  0]
 [ 0  0  0  0 22  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0 1400 0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  1  0  0  0  0  0  0  0  0  0]
 [ 2  0  0  0  0  0  0 1390 0  0  0  1  0  4]
 [ 0  0  0  0  0  0  0  2  2  0  0  0  0  0]
 [ 0  0  0  0  0  3  0  2  0 11  3  0  0  0]
 [ 0  0  0  0  0  0  0  2  0  0 13  0  0  0]
 [ 0  0  0  0  0  0  0  1  0  0  0 27  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0 10  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0 15]]
```

```
In [71]: print("Classification Report For K-Nearest Neighbors - 'Labels Test' & 'Prediction'")
print('-' * 97)
print(classification_report(y_test, k_nearest_neighbors_prediction))
```

Classification Report For K-Nearest Neighbors - 'Labels Test' & 'Prediction':

```
-----
              precision    recall  f1-score   support


     0           0.90       0.96       0.93         28
     1           0.00       0.00       0.00          1
     2           0.00       0.00       0.00          2
     3           0.00       0.00       0.00          1
     4           0.96       1.00       0.98         22
     6           1.00       1.00       1.00        1400
     7           0.00       0.00       0.00          1
     8           0.99       0.99       0.99        1397
     9           1.00       0.50       0.67          4
    10           1.00       0.58       0.73          19
    11           0.81       0.87       0.84          15
    12           0.96       0.96       0.96         28
    13           1.00       1.00       1.00         10
    14           0.79       1.00       0.88         15

 accuracy          0.99         2943
 macro avg         0.67         0.63         0.64         2943
 weighted avg      0.99         0.99         0.99         2943
```

```
In [72]: season = ["Accuracy", "Precision", "Recall", "F1-Score"]
season_score = [k_nearest_neighbors_accuracy, k_nearest_neighbors_precision,
                k_nearest_neighbors_recall, k_nearest_neighbors_f1_score]

sns.barplot(x=season, y=season_score, palette="rocket")
```

```
plt.title("K-Nearest Neighbors - Accuracy, Precision, Recall and F1-Score")
plt.xlabel("Evaluation Metrics")
plt.ylabel("Score")
plt.show()
```



Ada Boost

```
In [73]: # apply grid search to ada boost

from sklearn.ensemble import AdaBoostClassifier

ada_boost_model = AdaBoostClassifier()

ada_boost_param_grid = [
    {
        'n_estimators': [50, 100, 150, 200, 250, 300, 350, 400, 450, 500],
        'learning_rate': [0.001, 0.01, 0.1, 1.0, 10.0],
        'algorithm': ['SAMME', 'SAMME.R']
    }
]

ada_boost_grid_search = GridSearchCV(ada_boost_model, ada_boost_param_grid)
ada_boost_grid_search.fit(X_train, y_train)
ada_boost_grid_search.best_params_
```

```
Out[73]: {'algorithm': 'SAMME', 'learning_rate': 1.0, 'n_estimators': 100}
```

```
In [74]: print("Evaluation for Ada Boost".center(75, "_"))

ada_boast = AdaBoostClassifier(n_estimators=100, random_state=0)
ada_boast.fit(X_train, y_train)

ada_boast_prediction = ada_boast.predict(X_test)
ada_boast_accuracy = accuracy_score(y_test, ada_boast_prediction)
ada_boast_precision = precision_score(y_test, ada_boast_prediction, average="weight")
ada_boast_recall = recall_score(y_test, ada_boast_prediction, average="weighted")
ada_boast_f1_score = f1_score(y_test, ada_boast_prediction, average="weighted")

print("Prediciton:      ", ada_boast_prediction)
print("_" * 75)

print("Accuracy:" + "\t" + f"{{(ada_boast_accuracy * 100)}}%")
print("Precision:" + "\t" + f"{{(ada_boast_precision * 100)}}%")
print("Recall:" + "\t\t" + f"{{(ada_boast_recall * 100)}}%")
print("F1-Score:" + "\t" + f"{{(ada_boast_f1_score * 100)}}%")
print("_" * 75)
```

```
_____Evaluation for Ada Boost_____
Prediciton:      [6 8 6 ... 6 8 6]

Accuracy:      94.22358137954468%
Precision:      89.7850953107695%
Recall:      94.22358137954468%
F1-Score:      91.91177978962602%
```

```
In [75]: print("Confusion Matrix For Ada Boast - 'Labels Test' & 'Prediction':")
print('-' * 93)
print(confusion_matrix(y_test, ada_boast_prediction))
```

```
Confusion Matrix For Ada Boast - 'Labels Test' & 'Prediction':
```

```
-----
[[ 0  0  0  0  0  0  0  0  28  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  1]
 [ 0  0  0  0  0  0  0  0  0  0  0  2  0  0  0]
 [ 0  0  0  0  0  0  0  0  1  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  22  0  0  0  0  0  0  0]
 [ 0  0  0  0  0 1385  0  15  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  1  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  9  0 1388  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  4  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  3  0 16  0  0  0  0  0  0  0]
 [ 0  0  0  0  0 14  0  1  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0 28  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0 10  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0 15  0  0  0  0  0  0  0]]
```


```
In [76]: print("Classification Report For Ada Boast - 'Labels Test' & 'Prediction':")
print('-' * 97)
print(classification_report(y_test, ada_boast_prediction))
```

Classification Report For Ada Boost - 'Labels Test' & 'Prediction':

	precision	recall	f1-score	support
0	0.00	0.00	0.00	28
1	0.00	0.00	0.00	1
2	0.00	0.00	0.00	2
3	0.00	0.00	0.00	1
4	0.00	0.00	0.00	22
6	0.98	0.99	0.99	1400
7	0.00	0.00	0.00	1
8	0.91	0.99	0.95	1397
9	0.00	0.00	0.00	4
10	0.00	0.00	0.00	19
11	0.00	0.00	0.00	15
12	0.00	0.00	0.00	28
13	0.00	0.00	0.00	10
14	0.00	0.00	0.00	15
accuracy			0.94	2943
macro avg	0.13	0.14	0.14	2943
weighted avg	0.90	0.94	0.92	2943

```
In [77]: season = ["Accuracy", "Precision", "Recall", "F1-Score"]
season_score = [ada_boost_accuracy, ada_boost_precision,
                ada_boost_recall, ada_boost_f1_score]

sns.barplot(x=season, y=season_score, palette="rocket")
plt.title("Ada Boost - Accuracy, Precision, Recall and F1-Score")
plt.xlabel("Evaluation Metrics")
plt.ylabel("Score")
plt.show()
```



Gradient Boosting

In [102...

```
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV

gradient_boost_model = GradientBoostingClassifier(random_state=42)

gradient_boost_param_grid = {
    'learning_rate': [0.05, 0.1],
    'n_estimators': [100, 200],
    'max_depth': [3, 5],
    'criterion': ['friedman_mse']
}

gradient_boost_grid_search = GridSearchCV(
    estimator=gradient_boost_model,
    param_grid=gradient_boost_param_grid,
    cv=3,
    scoring='accuracy',
    n_jobs=-1
)

gradient_boost_grid_search.fit(X_train, y_train)

print(gradient_boost_grid_search.best_params_)
print(gradient_boost_grid_search.best_score_)
```

```
{'criterion': 'friedman_mse', 'learning_rate': 0.05, 'max_depth': 3, 'n_estimators': 200}
0.9960660513513596
```

```
In [103... print("Evaluation for Gradient Boosting".center(75, "_"))

gradient_boosting = GradientBoostingClassifier(
    n_estimators=100, learning_rate=1.0, max_depth=1, random_state=0
)
gradient_boosting.fit(X_train, y_train)

gradient_boosting_prediction = gradient_boosting.predict(X_test)
gradient_boosting_accuracy = accuracy_score(y_test, gradient_boosting_prediction)
gradient_boosting_precision = precision_score(
    y_test, gradient_boosting_prediction, average="weighted"
)
gradient_boosting_recall = recall_score(
    y_test, gradient_boosting_prediction, average="weighted"
)
gradient_boosting_f1_score = f1_score(
    y_test, gradient_boosting_prediction, average="weighted"
)

print("Predicition: ", gradient_boosting_prediction)
print("_" * 75)

print("Accuracy:" + "\t" + f"{(gradient_boosting_accuracy * 100)}%")
print("Precision:" + "\t" + f"{(gradient_boosting_precision * 100)}%")
print("Recall:" + "\t\t" + f"{(gradient_boosting_recall * 100)}%")
print("F1-Score:" + "\t" + f"{(gradient_boosting_f1_score * 100)}%")
print("_" * 75)
```

Evaluation for Gradient Boosting	
Predicition:	[10 10 10 ... 10 10 10]
Accuracy:	1.4276002719238612%
Precision:	1.070717309145664%
Recall:	1.4276002719238612%
F1-Score:	1.1690131561714865%

```
In [104... print("Confusion Matrix For Gradient Boosting - 'Labels Test' & 'Prediction':")
print('-' * 93)
print(confusion_matrix(y_test, gradient_boosting_prediction))
```

Confusion Matrix For Gradient Boosting - 'Labels Test' & 'Prediction':

```
-----
[[ 0  0  0  0  0  0  0  0  0  0  25  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  1  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  1  0  0]
 [ 0  0  0 19  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0 1412  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  1  0  0]
 [ 0  0  0  2  0  0  0  0  0  0 1380  0  0]
 [ 0  0  0  1  0  0  0  0  0  0  0  3  0]
 [ 0  0  0  0  0  0  0  0  3  0  8  0  0]
 [ 0  0  0  0  0  0  0  0 18  0  1  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  4  0 29]
 [ 0  0  0  0  0  0  0  0  0  0  0 16  0]
 [ 0  0  0  0  0  0  0  0  0  0 18  0  0]]
```

```
In [105... print("Classification Report For Gradient Boosting - 'Labels Test' & 'Prediction':")
print('-' * 97)
print(classification_report(y_test, gradient_boosting_prediction))
```

Classification Report For Gradient Boosting - 'Labels Test' & 'Prediction':


```
-----
              precision    recall  f1-score   support

     0           0.00         0.00         0.00         25
     1           0.00         0.00         0.00          1
     2           0.00         0.00         0.00          1
     3           0.86         1.00         0.93         19
     4           0.00         0.00         0.00        1412
     5           0.00         0.00         0.00          1
     6           0.00         0.00         0.00       1382
     7           0.00         0.00         0.00          4
     8           0.14         0.27         0.19         11
     9           0.00         0.00         0.00         19
    10           0.00         0.12         0.00         33
    11           0.84         1.00         0.91         16
    12           0.00         0.00         0.00         18

 accuracy          0.01         0.01         0.01       2942
 macro avg         0.14         0.18         0.16       2942
 weighted avg      0.01         0.01         0.01       2942
```

```
In [106... season = ["Accuracy", "Precision", "Recall", "F1-Score"]
season_score = [gradient_boosting_accuracy, gradient_boosting_precision,
                gradient_boosting_recall, gradient_boosting_f1_score]

sns.barplot(x=season, y=season_score, palette="rocket")
plt.title("Gradient Boosting - Accuracy, Precision, Recall and F1-Score")
plt.xlabel("Evaluation Metrics")
plt.ylabel("Score")
plt.show()
```



Hist Gradient Boosting

```
In [84]: from sklearn.ensemble import HistGradientBoostingClassifier
from sklearn.model_selection import GridSearchCV

hist_model = HistGradientBoostingClassifier(random_state=42)

param_grid = {
    "learning_rate": [0.01, 0.1],
    "max_iter": [100, 200],
    "max_depth": [3, 5]
}

grid = GridSearchCV(
    hist_model,
    param_grid,
    cv=3,
    n_jobs=-1,
    scoring="accuracy"
)

grid.fit(X_train, y_train)

print(grid.best_params_)

{'learning_rate': 0.01, 'max_depth': 5, 'max_iter': 200}
```

```
In [85]: print("Evaluation for Hist Gradient Boosting".center(75, "_"))
```



```

hist_gradient_boosting = HistGradientBoostingClassifier(
    loss="log_loss", learning_rate=0.1, max_iter=100, max_leaf_nodes=31, random_state=42
)
hist_gradient_boosting.fit(X_train, y_train)

hist_gradient_boosting_prediction = hist_gradient_boosting.predict(X_test)
hist_gradient_boosting_accuracy = accuracy_score(
    y_test, hist_gradient_boosting_prediction
)
hist_gradient_boosting_precision = precision_score(
    y_test, hist_gradient_boosting_prediction, average="weighted"
)
hist_gradient_boosting_recall = recall_score(
    y_test, hist_gradient_boosting_prediction, average="weighted"
)
hist_gradient_boosting_f1_score = f1_score(
    y_test, hist_gradient_boosting_prediction, average="weighted"
)

print("Predictions: ", hist_gradient_boosting_prediction)
print("_" * 75)

print("Accuracy:" + "\t" + f"{(hist_gradient_boosting_accuracy * 100)}%")
print("Precision:" + "\t" + f"{(hist_gradient_boosting_precision * 100)}%")
print("Recall:" + "\t\t" + f"{(hist_gradient_boosting_recall * 100)}%")
print("F1-Score:" + "\t" + f"{(hist_gradient_boosting_f1_score * 100)}%")
print("_" * 75)

```

Evaluation for Hist Gradient Boosting	
Predictions:	[6 8 6 ... 6 6 6]
Accuracy:	83.11247026843357%
Precision:	80.6064688921018%
Recall:	83.11247026843357%
F1-Score:	81.4489140901371%

```

In [86]: print("Confusion Matrix For Hist Gradient Boosting - 'Labels Test' & 'Prediction':")
print('-' * 93)
print(confusion_matrix(y_test, hist_gradient_boosting_prediction))


```

Confusion Matrix For Hist Gradient Boosting - 'Labels Test' & 'Prediction':

```
-----
[[ 0  0  0  0  0  0  0  0 28  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  1  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  2  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  1  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0 22  0  0  0  0  0  0]
 [ 0  0  0  0  10 1335  0 35  0  1  4  0 14  1]
 [ 0  0  0  0  0  0  0  1  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0 298  0 1099  0  0  0  0  0  0]
 [ 0  0  0  0  0  1  0  3  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  3  0 16  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  4  0  0  0  0 11  0  0  0  0]
 [ 0  0  0  0  0  3  0 25  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  8  0  1  0  0  0  0  1  0  0]
 [ 0  0  0  0  0  0  0 15  0  0  0  0  0  0  0]]
```

```
In [87]: season = ["Accuracy", "Precision", "Recall", "F1-Score"]
season_score = [
    hist_gradient_boosting_accuracy,
    hist_gradient_boosting_precision,
    hist_gradient_boosting_recall,
    hist_gradient_boosting_f1_score,
]

sns.barplot(x=season, y=season_score, palette="rocket")
plt.title("Hist Gradient Boosting - Accuracy, Precision, Recall and F1-Score")
plt.xlabel("Evaluation Metrics")
plt.ylabel("Score")
plt.show()
```



XGBoost Classifier

```
In [88]: import pandas as pd

# Count samples in each class
counts = dataset["connection"].value_counts()

# Keep only classes having at least 2 samples
valid_classes = counts[counts >= 2].index

dataset = dataset[dataset["connection"].isin(valid_classes)].reset_index(drop=True)

print(dataset.shape)
print(dataset["connection"].value_counts())
```

```
(9805, 42)
connection
6      4706
8      4605
12     111
0       82
4       63
11      63
14      59
13      53
10      38
9       13
7        5
2        4
1        3
Name: count, dtype: int64
```

```
In [89]: X = dataset.drop("connection", axis=1)
        y = dataset["connection"]
```

```
In [90]: from sklearn.preprocessing import LabelEncoder

        obj_cols = X.select_dtypes(include="object").columns

        for col in obj_cols:
            le = LabelEncoder()
            X[col] = le.fit_transform(X[col])
```

```
In [91]: target_encoder = LabelEncoder()
        y = target_encoder.fit_transform(y)

        print("Classes:", target_encoder.classes_)
        print("Encoded labels:", sorted(set(y)))
```

```
Classes: [ 0  1  2  4  6  7  8  9 10 11 12 13 14]
Encoded labels: [np.int64(0), np.int64(1), np.int64(2), np.int64(3), np.int64(4), n
p.int64(5), np.int64(6), np.int64(7), np.int64(8), np.int64(9), np.int64(10), np.int
64(11), np.int64(12)]
```

```
In [92]: print(X.shape)
        print(len(y))
```

```
(9805, 41)
9805
```

```
In [93]: from sklearn.preprocessing import StandardScaler

        scaler = StandardScaler()
        X_scaled = scaler.fit_transform(X)

        print(X_scaled.shape)
```

```
(9805, 41)
```

```
In [94]: from sklearn.model_selection import train_test_split

        X_train, X_test, y_train, y_test = train_test_split(
```

```

X_scaled,
y,
test_size=0.3,
random_state=8,
stratify=y
)

print(np.unique(y_train))
print(np.unique(y_test))

```

```

[ 0  1  2  3  4  5  6  7  8  9 10 11 12]
[ 0  1  2  3  4  5  6  7  8  9 10 11 12]

```

In [95]: `from xgboost import XGBClassifier`

```

model = XGBClassifier(
    objective="multi:softprob",
    eval_metric="mlogloss",
    random_state=42
)

model.fit(X_train, y_train)

```

Out[95]:

```

XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds
               =None,
               enable_categorical=True, eval_metric='mlogloss',
               feature_types=None, feature_weights=None, gamma=None,
               grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin
               =None,

```

In [96]: `print("Evaluation for XGBoost".center(75, '_'))`

```

xgboost = XGBClassifier()
xgboost.fit(X_train, y_train)

xgboost_prediction = xgboost.predict(X_test)
xgboost_accuracy = accuracy_score(y_test, xgboost_prediction)
xgboost_precision = precision_score(
    y_test, xgboost_prediction, average="weighted")
xgboost_recall = recall_score(y_test, xgboost_prediction, average="weighted")
xgboost_f1_score = f1_score(y_test, xgboost_prediction, average="weighted")

print("Predictions: ", xgboost_prediction)
print('_' * 75)

print("Accuracy:" + "\t" + f"{(xgboost_accuracy * 100)}%")
print("Precision:" + "\t" + f"{(xgboost_precision * 100)}%")

```

```
print("Recall:" + "\t\t" + f"{(xgboost_recall * 100)}%")
print("F1-Score:" + "\t" + f"{(xgboost_f1_score * 100)}%")
```

Evaluation for XGBoost

Prediction: [6 6 6 ... 6 4 4]

Accuracy: 99.8640380693406%
 Precision: 99.84879150685241%
 Recall: 99.8640380693406%
 F1-Score: 99.8511626833975%

```
In [97]: print("Confusion Matrix For XGBoost - 'Labels Test' & 'Prediction':")
print('-' * 93)
print(confusion_matrix(y_test, xgboost_prediction))
```

Confusion Matrix For XGBoost - 'Labels Test' & 'Prediction':

```
-----
-----
[[ 25  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  1  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  1  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0 19  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0 1412 0  0  0  0  0  0  0  0]
 [  0  0  0  1  0  0  0  0  0  0  0  0  0]
 [  0  1  0  0  0  0 1381 0  0  0  0  0  0]
 [  0  0  0  0  0  0  0  4  0  0  0  0  0]
 [  0  0  0  0  0  0  0  0 11  0  0  0  0]
 [  0  0  0  0  0  0  0  0  0 19  0  0  0]
 [  0  0  0  0  0  0  0  0  0  0 33  0  0]
 [  0  0  0  0  0  0  0  0  0  0  0 16  0]
 [  0  0  0  0  0  0  2  0  0  0  0  0 16]]
```


```
In [98]: print("Classification Report For XGBoost - 'Labels Test' & 'Prediction':")
print("-" * 97)
print(classification_report(y_test, xgboost_prediction))
```

Classification Report For XGBoost - 'Labels Test' & 'Prediction':

	precision	recall	f1-score	support
0	1.00	1.00	1.00	25
1	0.50	1.00	0.67	1
2	1.00	1.00	1.00	1
3	0.95	1.00	0.97	19
4	1.00	1.00	1.00	1412
5	0.00	0.00	0.00	1
6	1.00	1.00	1.00	1382
7	1.00	1.00	1.00	4
8	1.00	1.00	1.00	11
9	1.00	1.00	1.00	19
10	1.00	1.00	1.00	33
11	1.00	1.00	1.00	16
12	1.00	0.89	0.94	18
accuracy			1.00	2942
macro avg	0.88	0.91	0.89	2942
weighted avg	1.00	1.00	1.00	2942

```
In [99]: season = ["Accuracy", "Precision", "Recall", "F1-Score"]
season_score = [xgboost_accuracy, xgboost_precision, xgboost_recall, xgboost_f1_sco

sns.barplot(x=season, y=season_score, palette="rocket")
plt.title("XGBoost - Accuracy, Precision, Recall and F1-Score")
plt.xlabel("Evaluation Metrics")
plt.ylabel("Score")
plt.show()
```



Comparison

In [108...

```
import matplotlib.pyplot as plt
import numpy as np

models = [
    "KNN",
    "Decision Tree",
    "Random Forest",
    "Gradient Boosting",
    "Hist Gradient Boosting",
    "XGBoost",
    "AdaBoost",
]

accuracy = [
    k_nearest_neighbors_accuracy,
    decision_tree_accuracy,
    random_forest_classifier_accuracy,
    gradient_boosting_accuracy,
    hist_gradient_boosting_accuracy,
    xgboost_accuracy,
    ada_boost_accuracy,
]

precision = [
    k_nearest_neighbors_precision,
    decision_tree_precision,
    random_forest_classifier_precision,
    gradient_boosting_precision,
```



```

    hist_gradient_boosting_precision,
    xgboost_precision,
    ada_boost_precision,
]
recall = [
    k_nearest_neighbors_recall,
    decision_tree_recall,
    random_forest_classifier_recall,
    gradient_boosting_recall,
    hist_gradient_boosting_recall,
    xgboost_recall,
    ada_boost_recall,
]
f1_score = [
    k_nearest_neighbors_f1_score,
    decision_tree_f1_score,
    random_forest_classifier_f1_score,
    gradient_boosting_f1_score,
    hist_gradient_boosting_f1_score,
    xgboost_f1_score,
    ada_boost_f1_score,
]

barWidth = 0.15

r1 = np.arange(len(accuracy))
r2 = [x + barWidth for x in r1]
r3 = [x + barWidth for x in r2]
r4 = [x + barWidth for x in r3]


plt.figure(figsize=(15, 5))

plt.bar(r1, accuracy, color="red", width=barWidth, edgecolor="white", label="Accuracy")
plt.bar(
    r2, precision, color="green", width=barWidth, edgecolor="white", label="Precision"
)
plt.bar(r3, recall, color="blue", width=barWidth, edgecolor="white", label="Recall")
plt.bar(
    r4, f1_score, color="yellow", width=barWidth, edgecolor="white", label="F1-Score"
)

plt.xlabel("Models", fontweight="bold")
plt.xticks([r + barWidth for r in range(len(accuracy))], models)

plt.legend(bbox_to_anchor=(1.05, 1), loc="upper left", borderaxespad=0.0)
plt.show()

```



Deep Learning

```
In [110...] import numpy as np

num_classes = len(np.unique(y_train))
print("Number of classes:", num_classes)
```

Number of classes: 13

```
In [111...] X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], 1)
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], 1)
```

```
In [112...] print(np.unique(y_train))
```

[0 1 2 3 4 5 6 7 8 9 10 11 12]

```
In [113...] import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten, Dense

# Number of input features
num_features = X_train.shape[1]

# Number of output classes
num_classes = len(np.unique(y_train))

# Reshape for CNN
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], 1)
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], 1)

# Build model
model = Sequential([
    Conv1D(32, kernel_size=3, activation='relu',
           input_shape=(num_features, 1)),
    MaxPooling1D(pool_size=2),
    Flatten(),
    Dense(64, activation='relu'),
    Dense(num_classes, activation='softmax')
])

# Compile
```

```

model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])

# Train
model.fit(X_train, y_train, epochs=10, batch_size=32)

# Evaluate
loss, accuracy = model.evaluate(X_test, y_test)
print("Accuracy:", accuracy)

```

WARNING:tensorflow:TensorFlow GPU support is not available on native Windows for TensorFlow >= 2.11. Even if CUDA/cuDNN are installed, GPU will not be used. Please use WSL2 or the TensorFlow-DirectML plugin.

Epoch 1/10

215/215 ————— 2s 3ms/step - accuracy: 0.9672 - loss: 0.2474

Epoch 2/10

215/215 ————— 1s 2ms/step - accuracy: 0.9889 - loss: 0.0486

Epoch 3/10

215/215 ————— 1s 2ms/step - accuracy: 0.9904 - loss: 0.0378

Epoch 4/10

215/215 ————— 1s 2ms/step - accuracy: 0.9932 - loss: 0.0292

Epoch 5/10

215/215 ————— 1s 2ms/step - accuracy: 0.9939 - loss: 0.0230

Epoch 6/10

215/215 ————— 1s 2ms/step - accuracy: 0.9946 - loss: 0.0201

Epoch 7/10

215/215 ————— 1s 2ms/step - accuracy: 0.9955 - loss: 0.0168

Epoch 8/10

215/215 ————— 1s 2ms/step - accuracy: 0.9958 - loss: 0.0160

Epoch 9/10

215/215 ————— 1s 3ms/step - accuracy: 0.9966 - loss: 0.0131

Epoch 10/10

215/215 ————— 1s 3ms/step - accuracy: 0.9966 - loss: 0.0132

92/92 ————— 0s 2ms/step - accuracy: 0.9956 - loss: 0.0248

Accuracy: 0.9955812096595764

```

In [115... print("Prediction:", model)
            print("_" * 75)

            print(f"Accuracy:\t{accuracy*100:.2f}%")
            print(f"Loss:\t\t{loss:.4f}")

```

Prediction: <Sequential name=sequential_1, built=True>

Accuracy: 99.56%

Loss: 0.0248


```

In [118... season = ["Accuracy", "Loss"]
            season_score = [accuracy, loss]

            sns.barplot(x=season, y=season_score, palette="rocket")
            plt.title("CNN - Accuracy and Loss")
            plt.xlabel("Evaluation Metrics")

```

```
plt.ylabel("Score")
plt.show()
```



Autoencoder

```
In [119... import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

# Define the Autoencoder model
model = Sequential()
model.add(Dense(128, activation="relu", input_shape=(num_features,)))
model.add(Dense(64, activation="relu"))
model.add(Dense(128, activation="relu"))
model.add(Dense(num_features, activation="linear"))

# Compile the model
model.compile(optimizer="adam", loss="mean_squared_error")

# Train the model
model.fit(X_train, X_train, epochs=10, batch_size=32)

# Use the trained model for anomaly detection
reconstructed_data = model.predict(X_test)
auto_encoder_loss = tf.keras.losses.mae(reconstructed_data, X_test)
auto_encoder_accuracy = tf.keras.metrics.binary_accuracy(reconstructed_data, X_test)
```

```

Epoch 1/10
215/215 ————— 2s 2ms/step - loss: 0.3094
Epoch 2/10
215/215 ————— 1s 3ms/step - loss: 0.0968
Epoch 3/10
215/215 ————— 1s 3ms/step - loss: 0.0486
Epoch 4/10
215/215 ————— 1s 2ms/step - loss: 0.0386
Epoch 5/10
215/215 ————— 1s 2ms/step - loss: 0.0421
Epoch 6/10
215/215 ————— 1s 2ms/step - loss: 0.0386
Epoch 7/10
215/215 ————— 1s 2ms/step - loss: 0.0231
Epoch 8/10
215/215 ————— 1s 3ms/step - loss: 0.0180
Epoch 9/10
215/215 ————— 1s 3ms/step - loss: 0.0228
Epoch 10/10
215/215 ————— 1s 3ms/step - loss: 0.0114
92/92 ————— 0s 2ms/step

```

In [120...

```

print("Predicition: ", model)
print('_' * 75)

print("Accuracy:" + "\t" + f"{(auto_encoder_accuracy * 100)}%")
print("Loss:" + "\t\t" + f"{(auto_encoder_loss * 100)}%")

```

Predicition: <Sequential name=sequential_2, built=True>

```

Accuracy:      [[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]]%
Loss:          [ 5.05728337 10.79346112  5.75157837 ...  9.66396504  2.96880651
 2.87822754]%


```

In [121...

```

# plot the loss distribution of the test set
plt.hist(auto_encoder_loss, bins=50)
plt.xlabel("Autoencoder Loss")
plt.ylabel("No of examples")
plt.show()

```



Gated Recurrent Unit (GRU)

In [122...

```
# Apply gru to the dataset
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import GRU, Dense

# Define the GRU model
model = Sequential()
model.add(GRU(32, input_shape=(num_features, 1)))
model.add(Dense(num_classes, activation="softmax"))

# Compile the model
model.compile(
    optimizer="adam", loss="sparse_categorical_crossentropy", metrics=["accuracy"]
)

# Train the model
model.fit(X_train, y_train, epochs=10, batch_size=32)

# Evaluate the model
gru_loss, gru_accuracy = model.evaluate(X_test, y_test)
```

```

Epoch 1/10
215/215 ————— 9s 24ms/step - accuracy: 0.8375 - loss: 0.9464
Epoch 2/10
215/215 ————— 5s 22ms/step - accuracy: 0.9325 - loss: 0.2773
Epoch 3/10
215/215 ————— 5s 23ms/step - accuracy: 0.9491 - loss: 0.2028
Epoch 4/10
215/215 ————— 5s 25ms/step - accuracy: 0.9643 - loss: 0.1722
Epoch 5/10
215/215 ————— 5s 25ms/step - accuracy: 0.9668 - loss: 0.1481
Epoch 6/10
215/215 ————— 5s 24ms/step - accuracy: 0.9675 - loss: 0.1321
Epoch 7/10
215/215 ————— 5s 23ms/step - accuracy: 0.9704 - loss: 0.1189
Epoch 8/10
215/215 ————— 5s 22ms/step - accuracy: 0.9765 - loss: 0.1053
Epoch 9/10
215/215 ————— 4s 21ms/step - accuracy: 0.9787 - loss: 0.0953
Epoch 10/10
215/215 ————— 4s 20ms/step - accuracy: 0.9786 - loss: 0.0901
92/92 ————— 1s 9ms/step - accuracy: 0.9782 - loss: 0.0883

```

```

In [123... print("Predicition: ", model)
            print('_' * 75)

            print("Accuracy:" + "\t" + f"{(gru_accuracy * 100)}%")
            print("Loss:" + "\t\t" + f"{(gru_loss * 100)}%")

Predicition:      <Sequential name=sequential_3, built=True>

```

```

Accuracy:      97.82460927963257%
Loss:          8.827356994152069%

```

```

In [124... season = ["Accuracy", "Loss"]
            season_score = [gru_accuracy, gru_loss]

            sns.barplot(x=season, y=season_score, palette="rocket")
            plt.title("CNN - Accuracy and Loss")
            plt.xlabel("Evaluation Metrics")
            plt.ylabel("Score")
            plt.show()

```

